



Exploratory Data Analysis in R

Session 6
Amrom Obstfeld
Introduction to R
Workshop
May 6, 2019


Goals and Objectives

- Appreciate the role and process of exploratory data analysis (EDA) in understanding data
- Learn about MIMIC-III data set and its utility in learning to work with biomedical data
- Further reinforce the skills learned during the previous five sessions



What is EDA

Exploratory Data Analysis



A
systematic
informal
process to
learn about
your data

Exploratory Data Analysis

EDA is an iterative process in which we:

1. Generate questions about our data.
2. Search for answers by visualizing, transforming, and modelling our data.
3. Use what we learn to refine your questions and/or generate new questions.

Exploratory Data Analysis

Two Fundamental Questions


1. What is the distribution of data in each of my variables
2. How do my variables relate with one another

How do we explore variation in data?

Quantitative

Range, mean,
median, mode


Histograms





Categorical

Frequency (counts, percent)

Bar chart



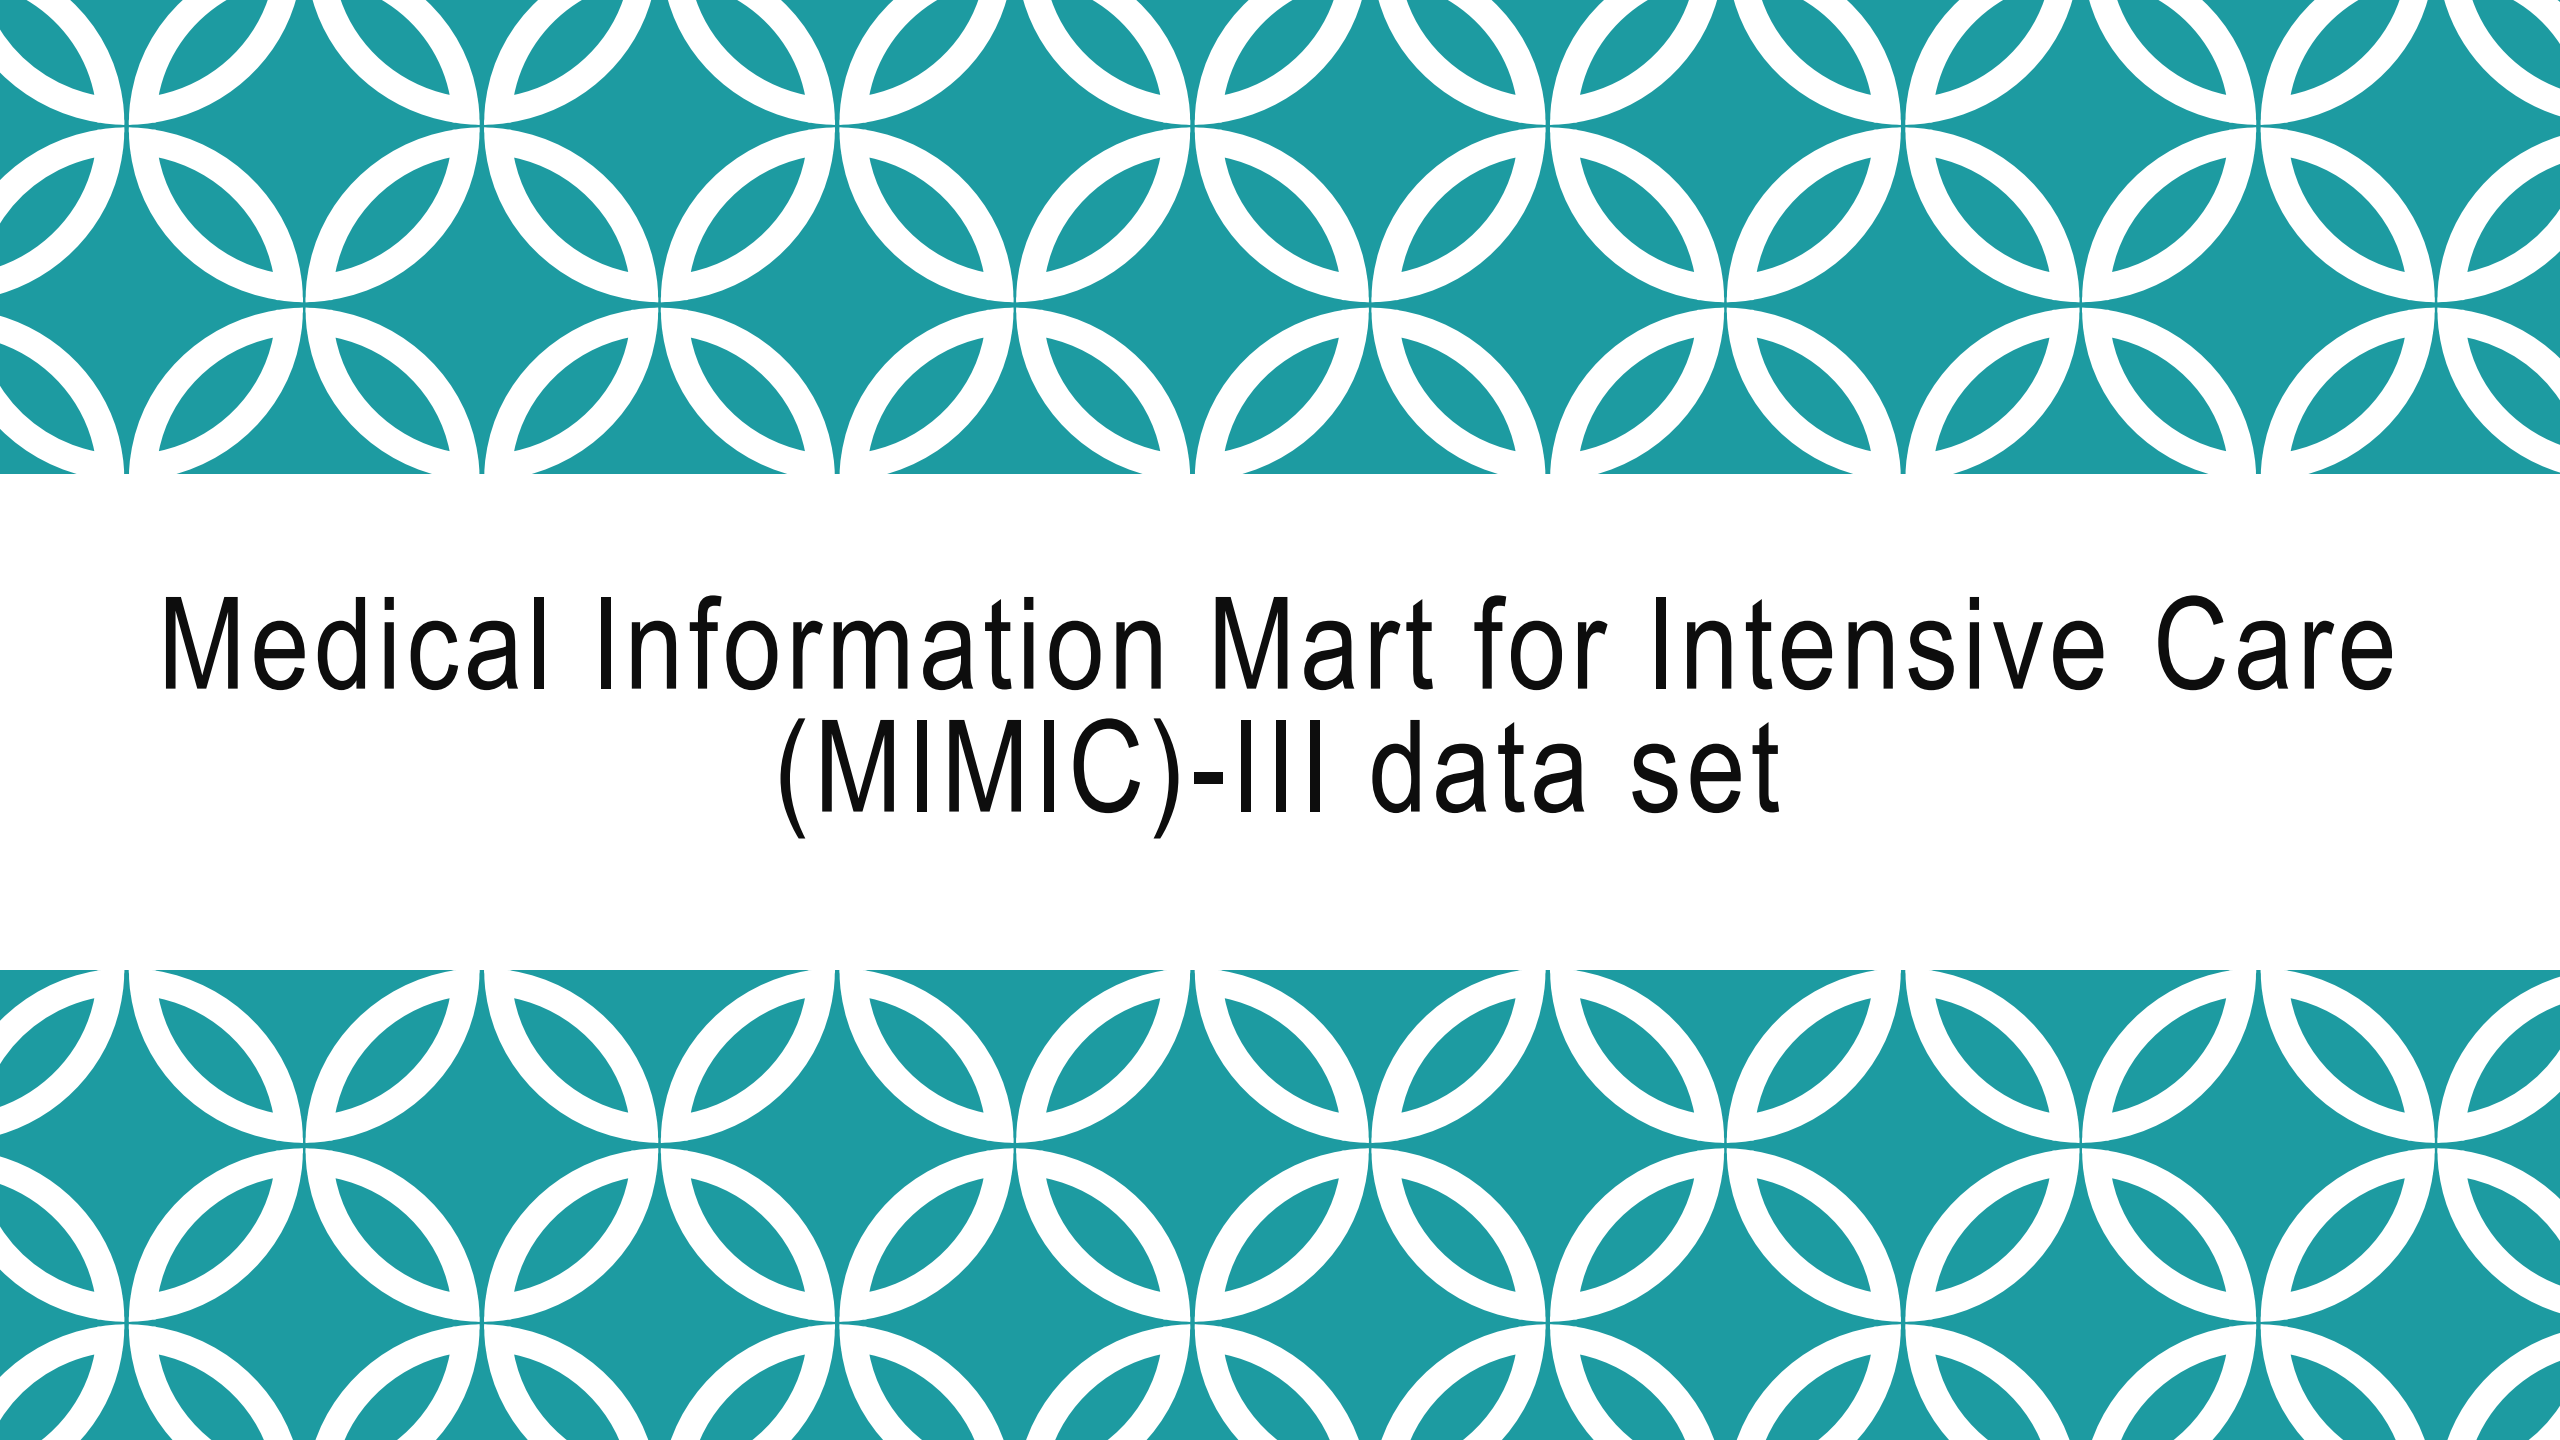
How do we explore COvariation in data?



EDA as Data QC

Investigate:


- Missing data
- Outliers
- Logical discrepancies



Medical Information Mart for Intensive Care (MIMIC)-III data set

MIMIC Data set

- MIMIC-III is a widely-used, freely available biomedical dataset
- Developed by the MIT Lab for Computational Physiology
- Deidentified health data associated with >40,000 critical care patients
- Includes demographics, vital signs, laboratory tests, medications, free text notes and more
- Details are available at <https://mimic.physionet.org/>



MIMIC Data Access

- Course in protecting human research participants including HIPAA requirements.
- Data use agreement
 - outlines appropriate data usage
 - security standards
 - forbids efforts to identify individual patients

Common uses for MIMIC

- Educational coursework
 - medical analytics courses
- Research
 - machine learning approaches for prediction of patient outcomes
 - semantic analysis of unstructured patient note
- Datathons



EDA Exercise



EDA exercise outline

For each exercise:

- Explicit task description
- Use functions to complete the task
- Question-set based on your findings

EDA exercise outline



Your Turn

Introduce yourself to your neighbors

- Who are you?
- Where are you from?
- What do you do with data?
- Have you ever used R?

3:00

```
1 ---
2 title: "R Notebook"
3 output: html_notebook
4 ---
5
6 This is an R Notebook. R Notebooks are written in R Markdown. An R
7 Notebook is like an electronic lab notebook, but for data analysis. You
8 can use R Notebooks to take notes, write code, and you can run that code
9 and see the results in the same document.
10
11 To take notes, simply edit the text in this document. For example, edit
12 the following line to replace XXX with your name:
13
14 My name is XXX, and I'm editing an R Notebook!
15
16 In an R Markdown document, code goes into code chunks. Each code chunk
17 starts with three back-ticks (```) and the letter "r" in curly brackets.
18 It ends with a line that only has three backticks (```). The RStudio
19 editor makes the background color of code chunks gray. This way it's easy
20 to see where all the code chunks are. You can run the code in a code
21 chunk by clicking the green triangle in the upper right corner of the
22 code chunk. The results will appear beneath the chunk. Try it!
23
24 ```{r}
25 plot(cars)
26 ```
27
28 Good job!
29
30 You can open a new R Notebook by going to File > New File > R
31 Notebook.
```



Exercise 0



Exercise 0

Use the code block below to run `library()` on
"tidyverse"

03:00

Exercise 1

Use the `read_csv()` function to read `mimic.csv` into a new data frame called "mimic"

Explore the data frame using some of the tools we've learned today (`summary()`, `head()`)

05:00

1. How many rows are in the data frame?
2. How many columns are in the data frame?
3. What does each row in the data frame represent?
4. How many columns contain information about the patient's admission and how many relate to the test order?

Exercise 2

- Use ``filter()`` to find the rows that have NA in the `valuenum` column
- Use ``select()`` to narrow down to just the "panel_test", "test_name", "component", "value" and "valuenum" columns
- Use ``arrange()`` to order the data frame by "value" and "component" columns

10:00

```
mimic %>%  
  filter(is.na(valuenum) %>%  
    select(panel_test, component,  
           value, valuenum) %>%  
  arrange(value, component)
```

1. What is the difference between the "value" and "valuenum" columns?
2. What kind of result values in the data set appear in the "value" column but are NA in the "valuenum" column?

Exercise 3

- Use `group_by()` and `summarize()` to get a sense for the counts of data in some of the columns with categorical data
- Use `n()` and `n_distinct` inside of the `summarize()` function to count the rows and distinct values in each categorical variable

10:00

```
mimic %>%  
  group_by(religion) %>%  
  summarise(d_pt = n_distinct(subject_id))
```

1. How many distinct patients and admissions are in the data frame?
2. How many different panel tests and components are in the data frame?
3. What is the most common religion for patients in the data frame?
4. *Challenge Question* What is the most commonly ordered test in the data set?

Exercise 4

- Use `ggplot()` to assess the distribution of `charttime` and `valuenum`
- Try different scales to visualize the laboratory results
- Use the `fill` aesthetic to parse out the contribution of different categorical variables, such as `category` or `fluid types`, to the distribution of laboratory values

10:00

```
ggplot(mimic)+  
  geom_histogram(aes(valuenum, fill=category))+  
  scale_x_log10()
```

1. What do you notice about the pattern of charttimes and valuenums in the data frame?
2. Are there differences between the distribution of results for "Chemistry" and "Hematology" test categorys?
3. Are there outlier categories with only a few results?
4. *Challenge Question* Can you estimate what the reference range is for the "Hemoglobin"?

Show and Tell 1 – What about NAs?

- Missing data (NA) is important to be aware of because its presence may indicate a problem with the data or may influence the statistics and conclusions we make

```
mimic %>%  
  mutate_all(is.na) %>%  
  summarise_all(sum) %>%  
  arrange(desc(value))
```


Show and Tell 2 – Characters in numeric data

- Analyzing laboratory data can be challenging for several reasons
- One reason is that numeric data can often be stored as characters

```
mimic %>%  
  filter(str_detect(value,">|<")) %>%  
  mutate(valuenum =  
    ifelse(str_detect(value,">|<"),  
           parse_number(value),  
           valuenum)) %>%  
  select(test_name, component, value,  
valuenum)
```

Show and Tell 3 – Working With Dates

- Medical data is inherently temporal in nature
- R has several packages that help in dealing with datetime based data

```
library(lubridate)
mimic %>%
  mutate(charthour = hour(charttime)) %>%
  filter(!category %in%
          c("HEMATOLOGY", "CHEMISTRY")) %>%
  distinct(curr_service, category, charttime,
            charthour, hadm_id, panel_test) %>%
  ggplot(aes(x=charthour, fill=category))+
  geom_bar()+
  facet_wrap(~curr_service, scales = "free_y")
```